

MNIST digit recognition

Josep Fortiana 2019-11-26

MNIST is a classic dataset, a standard testbed for classifying methods. In the introduction to the Kaggle competition we read: MNIST (“Modified National Institute of Standards and Technology”) is the de facto “Hello World” dataset of computer vision. Since its release in 1999, this classic dataset of handwritten images has served as the basis for benchmarking classification algorithms. As new machine learning techniques emerge, MNIST remains a reliable resource for researchers and learners alike.

The database has 60000 (*train*) and 10000 (*test*) 28×28 bitmaps of digits handwritten by high school students and employees of the US Census Bureau. Each image is labeled with the digit it represents.

The original data can be found at: Yann LeCun, Corinna Cortes, Christopher J.C. Burges, *The MNIST database of handwritten digits*, in a somewhat unusual format. A straightforward way to get it in a readily usable format is to load it from the *Keras/Tensorflow* framework. It is a rather large installation, but anyway we will need it for the Neural Networks lesson. We follow directions from JJ Allaire (2017) *Keras for R*. Alternatively, there is a .csv version in Joseph Chet Redmon’s web page, together with a python script to convert from the original format.

Warning! It is advisable to run the following three lines of code from RStudio or, still better, from the plain R console. It installs both the core *Keras* library as well as the *TensorFlow* backend.

```
#install.packages("keras",dependencies=TRUE,repos="https://cloud.r-project.org")
#library(keras)
#install_keras()
```

To obtain the MNIST data:

```
require(keras)
```

```
## Loading required package: keras
```

```
mnist <- dataset_mnist()
str(mnist)
```

```
## List of 2
## $ train:List of 2
## ..$ x: int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
## ..$ y: int [1:60000(1d)] 5 0 4 1 9 2 1 3 1 4 ...
## $ test :List of 2
## ..$ x: int [1:10000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
## ..$ y: int [1:10000(1d)] 7 2 1 0 4 1 4 9 5 9 ...
```

```
x_train <- mnist$train$x
y_train <- mnist$train$y
x_test <- mnist$test$x
y_test <- mnist$test$y
```

The x data is a 3-d array (images,width,height) of grayscale values. To prepare the data for training we convert the 3-d arrays into matrices by reshaping width and height into a single dimension (28×28 images are flattened into length 784 vectors). Then, we convert the grayscale values from integers ranging between 0 to 255 into floating point values ranging between 0 and 1:

```
# reshape
dim(x_train) <- c(nrow(x_train), 784)
dim(x_test) <- c(nrow(x_test), 784)
# rescale
x_train <- x_train / 255
x_test <- x_test / 255
```

The responses vectors `y_train` and `y_test` must be cast as factors to feed them to `rpart`

```
y_train<-as.factor(y_train)
y_test<-as.factor(y_test)
n_train<-length(y_train)
n_test<-length(y_test)
str(y_train)
```

```
## Factor w/ 10 levels "0","1","2","3",...: 6 1 5 2 10 3 2 4 2 5 ...
```

```
str(y_test)
```

```
## Factor w/ 10 levels "0","1","2","3",...: 8 3 2 1 5 2 5 10 6 10 ...
```

```
n_train
```

```
## [1] 60000
```

```
n_test
```

```
## [1] 10000
```

Some random images from the MNIST database:

```
k<-5
plot.zip.mosaic<-function(k){
  old.par<-par(mfrow=c(k,k),mar=c(1,1,1,1))
  Indexes<-sample(n_train,k^2)
  for (i in Indexes){
    m1<-t(matrix(x_train[i,],nrow=28,ncol=28))[,28:1]
    image(-m1,col=gray((0:255)/255),xaxt="n",yaxt="n")
  }
  par(old.par)}
options(repr.plot.width=5.5,repr.plot.height=5.5)
plot.zip.mosaic(k)
```

8	5	1	9	7
7	8	1	9	3
4	7	6	9	1
9	7	5	7	7
1	5	3	8	6